

Introduction to High Performance Computing

# **UIC HPC WORKSHOP**

## **Ramon Castillo**

HPC Systems Administrator, ACCC

## **Greg Cross**

HPC Technical Lead, ACCC

## **Himanshu Sharma**

HPC Project Manager, RRC

# Training Agenda

- Access
- Basic Linux Commands
- Storage on the Cluster
- Basic Script Creation
- Introduction to Job Submission
- Job Monitoring
- MPI Program Sample
- Additional Resources
- Questions & Answers
- Resources

# Access

## How to Access the Cluster

- Access requires a **UIC NetID** for login. Your password is your **ACCC Common Password**.
- Using Secure Shell (**SSH**) on Mac, Linux, and Unix:
  - Command requires terminal application window
  - **ssh netid@login-1.extreme.uic.edu**
- Using SSH on Windows
  - Requires third party application, such as **PuTTY**
  - Download PuTTY for free from <http://www.putty.org/>
    - From the PuTTY application select **SSH** radio button under **Session**
    - Enter **login-1.extreme.uic.edu** in **Host Name** field

# Access: File Transfers

- File transfers can be done via Secure Copy (**SCP**), which uses the same credentials as SSH.
- Using SCP on Mac, Linux, and Unix:
  - **scp -r *source\_path* *dest\_path***
  - Use **-r** when recursively copying contents of a directory
  - The remote path (either source or destination) is formatted as ***netid@login-1.extreme.uic.edu:/path***
- Using SCP on Windows:
  - PuTTY web site provides an analogous **PSCP** application which uses similar syntax to SCP.
  - **pscp.exe** is CLI-based; run within **cmd.exe** utility.

# Access: Useful SSH Flags and X11

## How to Access the Cluster

- To allow an **X11 GUI application** from within SSH:
  - `ssh -Y netid@login-1.extreme.uic.edu`
- The X11 server is already available on Linux and Unix workstations. **XQuartz** is freely available for the Mac; **Xming** is freely available for Windows.
- Using SSH in **verbose mode** (helpful when debugging):
  - `ssh -v netid@login-1.extreme.uic.edu`

# Basic Linux Commands: File System Navigation

- **cd**
  - The **change directory** command is used to navigate the file system:
    - **cd path**
- **ls**
  - The **list directory** command is used to gain information on files and subdirectories (folders) in a directory. A good example with useful flags:
    - **ls -lah**
    - **l** = long (detailed), **a** = all (including hidden dot files), **h** = human-readable file sizes
- **pwd**
  - The **print working directory** command is used to determine the directory in which you are currently located with respect to the file system hierarchy. This command is useful for determining where you are if you get lost navigating the file system.
- **mkdir**
  - By default, this command will **make** a new sub**directory** (folder) relative to the directory in which you are currently located. You may however provide an absolute path to where you wish a subdirectory to be made.
  - **mkdir test1** # creates test1 relative to current directory
  - **mkdir /export/home/netid/test2** # creates absolute path

# Basic Linux Commands: File I/O and Editing

- **cat**
  - outputs (and concatenates) the contents of one or more files to standard output (by default the screen)
  - **cat *filename filename ...***
- **grep**
  - Searches input globally for a regular expression and prints matching patterns (**regular expressions** are special character strings using a versatile search syntax)
  - **grep *regex filename***
- **vi** and **nano**
  - text-based file editors
  - **vi *filename* —or— nano *filename***
  - **nano** is more intuitive; **vi** is more powerful



# Basic Linux Commands: File Permissions

- **chmod** *mode filename filename ...*
  - changes **m**odes of read, write, and execute file permissions, as well as other persistent setuid and ownership (“sticky bit”).
  - Permissions are ordered by user, group, & other user access.
  - Current ownership can be determined using the long directory listing (**ls -l**); e.g., `-rwxr-x--` indicates the user can read, write, and execute, the group can read and execute, and world has no access (the leading `-` means no persistence flags are set).
  - Mode parameters can be relatively changed ***m±n*** (where ***m*** is one or more of u, g, or other; `+` to add or `-` to remove; and ***n*** is one or more of r, w, or x); e.g., **chmod ug+x filename** gives the user and group execute permission.
  - Mode parameters can be absolutely changed using octal-encoded values; e.g., **chmod 0750 filename** specifies `-rwxr-x--`.
  - N.b.: directories must have the execute bit set to be accessible.

# Basic Linux Commands: File Ownership

- **chown** *owner filename filename ...*
  - Changes the user (UIC NetID) owning the specified file(s) and/or directory(ies); wildcards such as **\*** can be used (with caution).
  - **owner** can be the user only or **user:group** (colon-separated).
- **chgrp** *group filename filename ...*
  - Changes just the group ownership of a file or a directory.
  - On Extreme, the group is either “**domain users**” or a lab share group name provided by support staff.

# Basic Linux Commands: Software & Modules

- **`which executable`**
  - This command shows the default path to the specified ***executable***. Extreme provides multiple versions of software packages through modules; the output of this command may change depending on the module(s) loaded. When testing interactively, use this command to verify you are running the version you want.
- **`module parameters ...`**
  - The **`module`** command is used to manage software packages on the cluster.
- **`module avail`**
  - Lists all the software modules available on the cluster.
- **`module load modulename`**
  - This command loads the software package into your path. Keep in mind you must use this command in your submit scripts in order to call software packages.
- **`module list`**
  - This command displays active modules listed in the order they were loaded.
- **`module unload modulename`**
  - This command removes the specified software package from your path.

# Basic Linux Commands: MPI Modules Example

There are multiple implementations of the Message Passing Interface (MPI), a standard parallel computing framework. The environment has to be consistent between the building and execution of an MPI application. In this example, note how the path to **mpirun** (an MPI execution harness) changes depending on the module loaded.

```
$ which mpirun
```

```
/opt/openmpi/bin/mpirun
```

```
$ module load compilers/intel
```

```
$ which mpirun
```

```
/export/share/compilers/intel/impi/4.1.1.036/intel64/bin/mpirun
```

```
$ module load tools/mpich-3.0.4-icc
```

```
$ which mpirun
```

```
/export/share/tools/mpich-3.0.4-icc/bin/mpirun
```

# Storage on the Cluster: File System Types & Paths

- **Home directories** are on **Persistent NFS Storage**
  - `/export/home/netid`
- ... as are **lab shares** (common space for research groups)
  - `/export/home/labshare`
- Fast temporary **scratch** filesystem is on **Lustre**
  - `/mnt/lustre/netid`

# Storage on the Cluster: Lab Share Specifics

- How to **request** a lab share
  - Have the head of your lab group send an email to [extreme@uic.edu](mailto:extreme@uic.edu) specifying the need, beginning membership, the lab's cluster affiliation (department- or faculty-funded), and the point of contact for the lab/project responsible for maintaining the member access list.
- How to **gain access** to a lab share that already exists.
  - Send an email to [extreme@uic.edu](mailto:extreme@uic.edu) requesting access and cc: the faculty member or their designated membership manager. Please notify the responsible party in advance to ensure they receive and approve your request.

# Basic Job Scripting: Parameter Checklist

- What are the main things to confirm when submitting a job?
  - **Cores** (number of CPU cores × number of nodes)
  - **Memory**
  - **Wallclock time** (expected duration of job run)
  - What is the **working directory** of your job
  - How to specify **input and output files**
  - What happens when you do not specify one of these parameters

# Basic Job Scripting: Example Submit Script

```
#!/bin/bash
#PBS -l mem=20gb
#PBS -l walltime=20:00:00
#PBS -l nodes=1:ppn=8
#PBS -j oe
#PBS -m abe
#PBS -M email_address
#PBS -N jobname
#PBS -d /export/home/netid/work_dir

sleep 30
```



# Basic Job Scripting: #PBS Headers

```
#!/bin/bash
```

Always specify the shell that the job script uses.

```
#PBS -l mem=20gb
```

This optional line tells the cluster how much memory your job intends to use and ensures that there is enough memory on the assigned nodes when you submit.

```
#PBS -l walltime=20:00:00
```

This line tells the cluster how long the job should run (HH:MM:SS).

```
#PBS -l nodes=1:ppn=8
```

This line specifies the number of nodes (physical compute nodes) and then the number of cores (processors) the job will need to run. The assigned resource is a product of these values.

# Basic Job Scripting: #PBS Headers (con't.)

**#PBS -j oe**

Allows the user to join and otherwise manipulate the standard output and standard error into a single file.

**#PBS -m a**b**e**

Requests a status email when a job begins, ends, or aborts.

**#PBS -M *email\_address***

Provides an email address in conjunction with the status email flag above.

**#PBS -N *jobname***

Names the job with a custom label to allow a user to make it easily identifiable in the list of jobs (e.g., provided by **qstat**).

**#PBS -d */export/home/netid/work\_dir***

Specifies the initial working directory for your job, generally where your job's data and/or program reside.

# Basic Job Scripting: Running Your Application

```
sleep 30
```

Commands that follow the **#PBS** headers are what are actually executed on the cluster. In this example, the job simply sleeps for 30 seconds and then exits.

Depending on the complexity of your applications, you may also need to include **\$PBS\_...** variables. The most common is **\$PBS\_NODEFILE**, used in conjunction with MPI. This file contains one or more lines, each specifying a hostname for each core on the respective node(s) assigned to a job. Another useful variable is **\$PBS\_NP**, which is equal to the total number of cores assigned to the job (and is always equal to the line count of **\$PBS\_NODEFILE**).

# Introduction to Job Submission

- Submit the job using a submit script requesting **one node**:
  - `qsub -l nodes=1 submit_script`
- Now request all **20** cores (**p**rocessors **p**er **n**ode) on **two nodes**:
  - `qsub -l nodes=2:ppn=20 submit_script`

The `-l` (lowercase l) precedes a comma-separated resource list. `qsub` can accept multiple "`-l ...`" argument pairs. You may specify these parameters at the command line or in the script.

- Request an interactive job with `-I` (capital i):
  - `qsub -I -l nodes=2:ppn=16`

# Job Monitoring

- **showq**
  - Displays information on all jobs that are active, queued, or blocked.
  - To only display only your jobs, pipe the output through **grep**:  
`showq | grep netid`
- **qstat**
  - An alternative queue monitoring application.
- **checkjob**
  - Good for gaining detailed information on an individual job or determine why it failed to run.
    - **checkjob -v jobid**
      - Provides detailed information on the specified job and any error messages.
    - **checkjob -v -v jobid**
      - Provides not only detailed information on the specified job and error messages, but displays the output of your submitted script.
- **qdel**
  - To use this command to cancel one of your jobs , use:  
`qdel jobid`

# Basic Job Scripting: Invalid Resource Requests

~~#PBS -l mem=256gb~~

If you ask for a node with more than 128GB of memory the job will never run as each node only has 128 GB of RAM.

~~#PBS -l walltime=720:00:00~~

The maximum walltime is 240 hours (or 10 days). If you submit a job with a walltime longer than the maximum, the the job will not run.

~~#PBS -l nodes=1:ppn=128~~

The **ppn** value cannot exceed the per-compute core count. Extreme has 16-core “Generation 1” and 20-core “Generation 2” compute nodes. The default **batch** queue uses G1 nodes; **edu\_shared** uses G2; users should inquire about the node generation for other queues. A valid equivalent to a 1 × 128 configuration would be **nodes=8:ppn=16**.

# MPI Program Sample on Extreme

```
#include <mpi.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv) {

    int world_size;
    int rank;
    char hostname[256];
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;

    MPI_Init(&argc, &argv); // Initialize the MPI environment
    MPI_Comm_size(MPI_COMM_WORLD, &world_size); // get the total number of processes
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // get the processor rank number
    MPI_Get_processor_name(processor_name, &name_len); // get the processor name

    gethostname(hostname, 255); // non-MPI function to get the
hostname

    printf("Hello world! I am process number: %d from processor %s on host %s out of
%d processors\n", rank, processor_name, hostname, world_size);

    MPI_Finalize();

    return 0;

}
```

# Sample Script to Run MPI Program on Extreme

```
#!/bin/bash
#PBS -l nodes=2:ppn=20,walltime=1:00
#PBS -N MPIsample
#PBS -q edu_shared
#PBS -m abe
#PBS -M netid@uic.edu
#PBS -e mpitest.err
#PBS -o mpitest.out
#PBS -d /export/home/netid/MPIsample

module load compilers/intel

mpirun -machinefile $PBS_NODEFILE -np $PBS_NP ./mpitest
```



# Building and Running MPI (Example)

0. Copy (Recursively) the directory containing the sample code and job script to your home directory.

```
cp -R /export/share/classes/cs-ece566/MPIsample ~
```

1. Load Intel MPI module before compiling the program.

```
module load compilers/intel
```

2. Compile the program

```
cd ~/MPIsample
```

```
mpicc -o mpitest mpitest.c
```

3. Modify and submit the job script to the queue:

```
qsub mpitest.pbs
```

- **UIC Account Requests**

- <http://extreme.uic.edu/requests-policies/adding-a-new-userclass/>

- **Software Request**

- <http://extreme.uic.edu/requests-policies/software-request/>

- **Available Software**

- <http://extreme.uic.edu/knowledge-base/available-software/>

- **Service Alerts**

- <http://extreme.uic.edu/news-annoucements/service-alerts/>

# Additional Resources

- **General Resources**
  - <http://software-carpentry.org/lessons.html>
  - <https://www.coursera.org>
- **vi Editor Commands**
  - <http://www.lagmonster.org/docs/vi.html>
- **PBS Command Resource**
  - <https://www.osc.edu/supercomputing/batch-processing-at-osc/pbs-directives-summary>
  - [http://www.nas.nasa.gov/hecc/support/kb/Commonly-Used-PBS-Commands\\_174.html](http://www.nas.nasa.gov/hecc/support/kb/Commonly-Used-PBS-Commands_174.html)
- **PuTTY Tutorials**
  - [http://www.jfitz.com/tips/putty\\_config.html](http://www.jfitz.com/tips/putty_config.html)
  - <http://www.gamexe.net/other/beginner-guide-ssh/>

# Contact Us

- Questions, comments, and issues should be sent to [extreme@uic.edu](mailto:extreme@uic.edu)
- Additional resources can be found at <http://extreme.uic.edu>
- Follow **@UIC\_Extreme** on Twitter